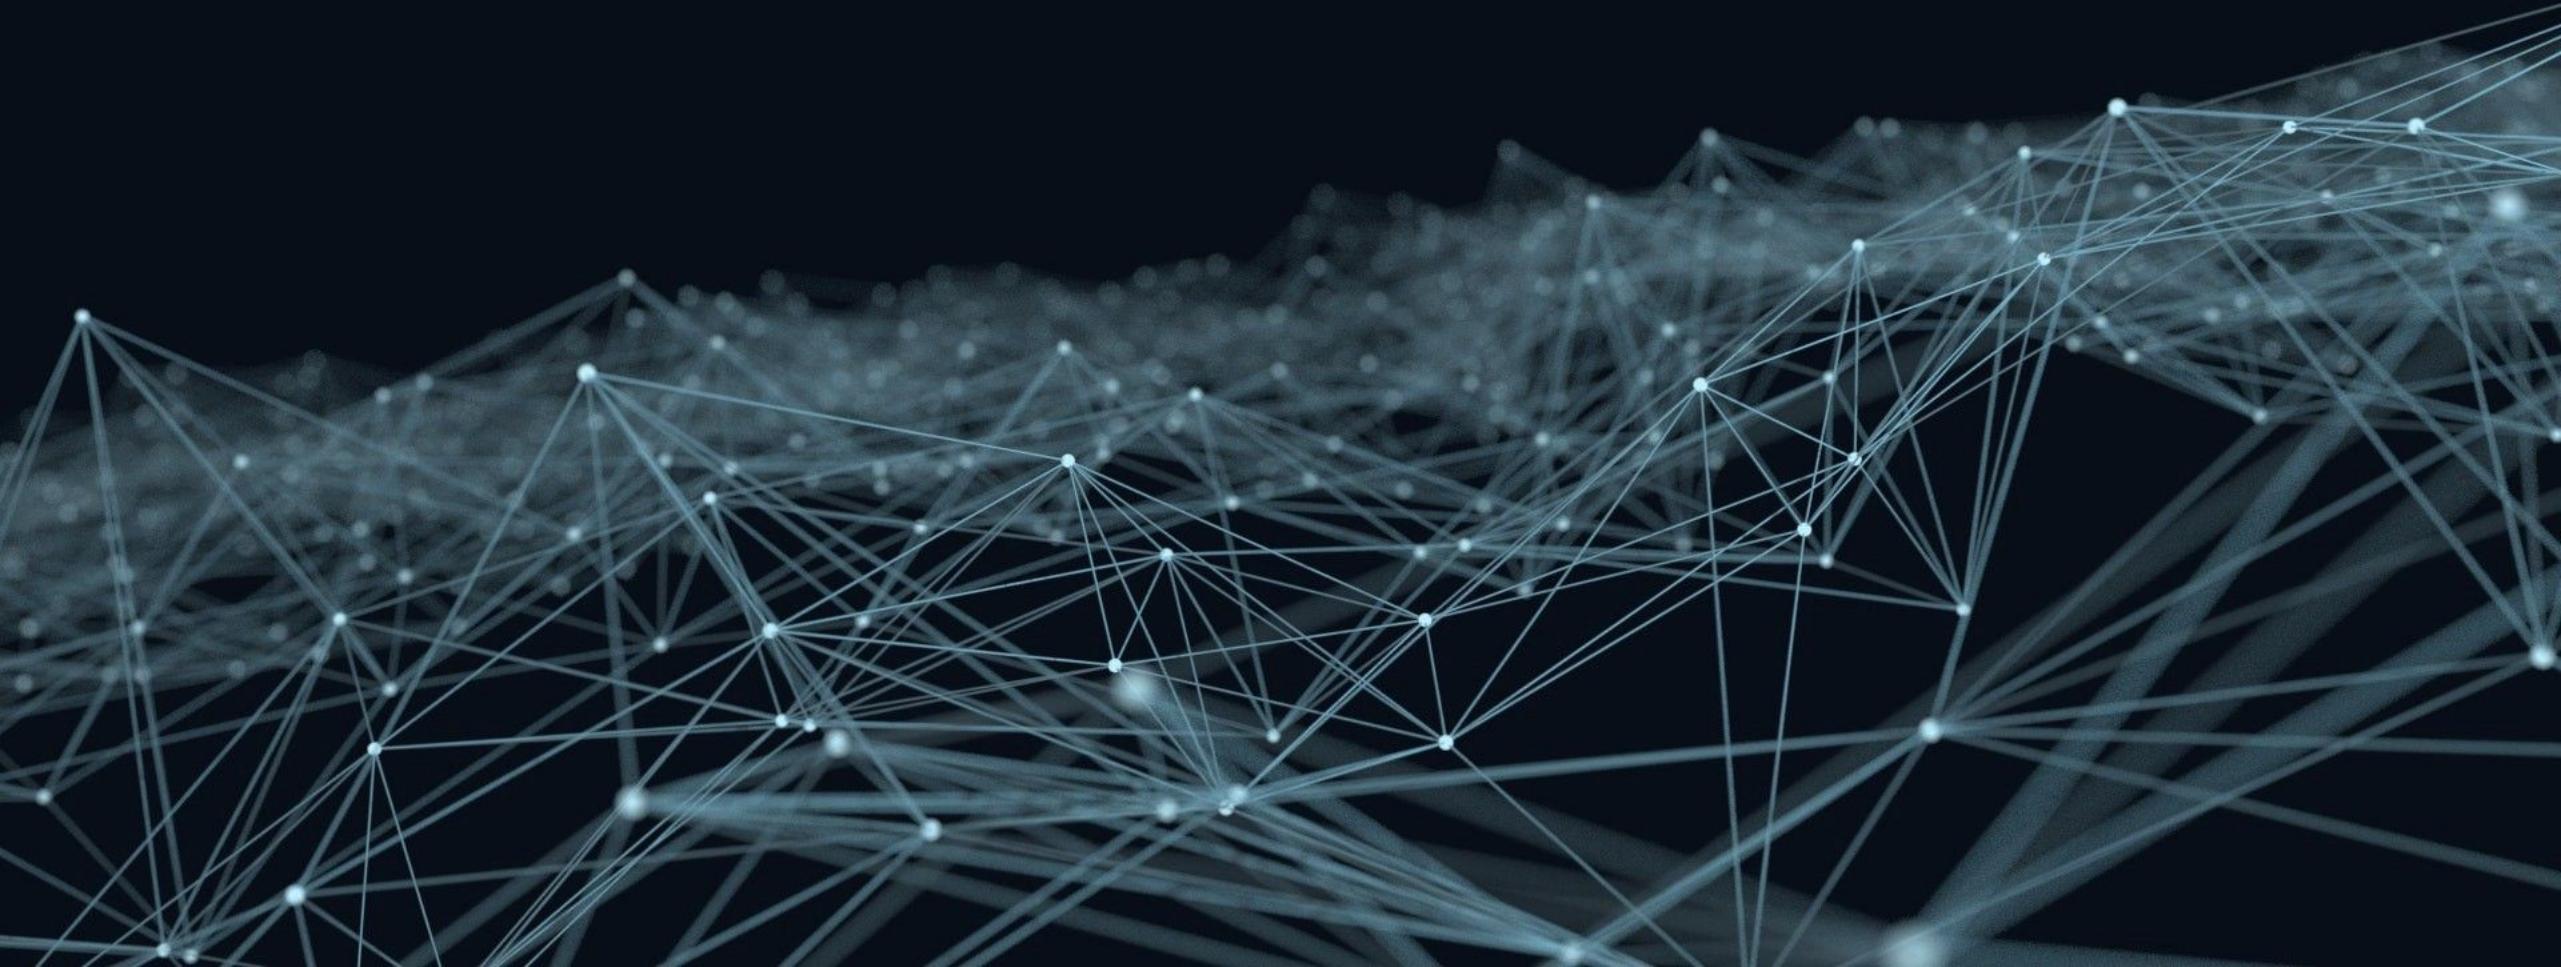


# Introduction to Jenkins

## Module 2: Jenkins Overview



# Module Topics

- Jenkins in modern DevOps ecosystems
- Jenkins LTS architecture and components
- Plugins, security model, and governance
- Controller/agent overview
- SCM-driven automation (GitHub, GitLab)
- Installing Jenkins LTS
- Running Jenkins in Docker
- Configuration-as-Code (JCasC)
- Initial setup wizard and global tools



# Origins

- In the early 2000s, most software teams
  - Integrated code infrequently
  - Ran builds manually
  - Discovered integration problems late in the development cycle
- In 2004 by Kohsuke Kawaguchi at Sun Microsystems created the Hudson project
  - Open-source written in Java
  - Designed to automate builds and tests
  - Integrated well with Java build tools like Ant and Maven
  - Easy to set up and extensible
  - Focused on automation rather than process enforcement



# Hudson to Jenkins (2011)

- In 2010–2011, Oracle acquired Sun Microsystems
  - This led to disagreements between the Hudson community and Oracle regarding
    - *Project governance*
    - *Trademark ownership*
    - *Contribution control*
- As a result
  - The community forked Hudson
  - The new project was named Jenkins
  - Most contributors and users moved to Jenkins
- From this point on
  - Jenkins became a community-driven project and the dominant tool for pipeline automation
  - Hudson gradually declined in relevance



# Job Centric CI (2011-2014)

- Early versions of Jenkins focused on
  - Freestyle jobs
  - UI-driven configuration
  - Scripted build steps
- Characteristics of early Jenkins
  - Configuration stored on the server
  - Heavy reliance on plugins
  - Tight coupling between built configurations and tool chains
  - Builds were often non-reproducible, especially after configuration changes
- This reflected the state of CI at the time
  - Automation existed, but IaC was not common, nor was the idea of a pipeline



# Pipeline as Code (2014-2016)

- As DevOps practices matured, teams demanded
  - Configuration of the pipeline logic needed to be versioned as the logic evolved
  - Reproducibility became a primary concern: consistent builds
  - Better support for complex workflows
- This led to the introduction of the Jenkins Pipeline and Jenkinsfile
- Innovations
  - Pipelines defined in code
  - Stored in Git repositories
  - Reviewed like application code
  - Support for stages, parallelism, and conditions
- This marked a major evolution in Jenkins' design philosophy



# DevOps Era (2016-Present)

- Jenkins evolved to support
  - Microservices
  - Containers and Docker
  - Kubernetes-based build agents
  - Cloud-native workflows
- Notable Trends
  - Declarative pipelines simplified syntax
  - Dynamic agents replaced static build servers
  - Integration with modern SCM platforms (GitHub, GitLab)
- Jenkins shifted from being “a CI tool” to a general-purpose automation engine



# Jenkins Today

- Why Jenkins is still relevant
  - Extreme flexibility
  - Large plugin ecosystem
  - Strong support for legacy and complex systems
- Challenges
  - Requires operational effort
  - Plugin compatibility management – many plugins are maintained by volunteers
  - Competition from managed CI/CD platforms like GitLab
- Despite newer tools, Jenkins remains widely used in
  - Large scale enterprises
  - Regulated environments
  - Complex, heterogeneous systems



# GitLab and GitHub

- GitHub Actions and GitLab CI/CD are modern alternatives to Jenkins pipelines
  - These are integrated pipeline tools built directly into their Git hosting services
- These repository tools treat CI/CD as a native feature of the repository
  - Jenkins = external automation engine
  - GitHub/GitLab = pipelines embedded in the SCM platform
- The same basic pipeline concepts apply to these as to Jenkins
  - The syntax and format of the configuration files differs
  - GitLab/GitHub use yaml configuration files
  - Jenkins uses groovy scripting



# GitHub Actions

- GitHub Actions
  - GitHub's built-in automation and CI/CD platform
  - Pipelines are defined using YAML workflow files stored in the repository
- Key characteristics
  - Tight integration with GitHub repositories
  - Event-driven (push, pull request, tag, issue events)
  - Pipelines live in the directory ".github/workflows"
  - Uses hosted runners or self-hosted runners
- Conceptual flow
  - GitHub event → Action workflow → Jobs → Steps
- Strengths
  - Zero infrastructure to manage
  - Simple setup for GitHub users
  - Strong marketplace of reusable actions



# GitLab CI/CD

- GitLab CI/CD is a first-class feature of GitLab, not an add-on.
  - Pipelines are defined in a file: “.gitlab-ci.yml”
- Key characteristics
  - CI/CD tightly integrated with GitLab repos
  - Built-in artifact storage and environment tracking
  - Native support for Docker and Kubernetes
  - GitLab runners execute jobs
- Conceptual flow
  - Git commit → GitLab pipeline → Stages → Jobs
- Strengths
  - Very cohesive developer experience
  - Strong DevOps lifecycle integration
  - Minimal configuration overhead



# Comparison

- Configuration style contrast
- Jenkins
  - Highly flexible
  - Scriptable pipelines (Groovy)
  - Powerful but verbose
  - Requires plugin management
- GitHub / GitLab
  - Opinionated YAML syntax
  - Simpler mental model
  - Less customization, faster onboarding
  - Platform constraints apply



# Jenkins vs GitHub Actions vs GitLab CI/CD

Feature	Jenkins	GitHub Actions	GitLab CI/CD
Hosting	Self-managed	SaaS	SaaS / Self
Pipeline Language	Jenkinsfile (Groovy)	YAML	YAML
Plugin Ecosystem	Very large	Marketplace actions	Built-in features
Customization	Extremely high	Moderate	High
Enterprise Control	Full	Limited	Strong

# Jenkins vs GitHub Actions vs GitLab CI/CD

- Jenkins is preferred for
  - Complex, multi-tool workflows
  - Hybrid or on-prem environments
  - Advanced customization needs
  - Regulated industries
  - Long-running or specialized builds
- GitHub/GitLab are preferred for
  - Simple to moderate pipelines
  - Tight integration with SCM
  - Minimal infrastructure management
  - Fast onboarding



# Comparison

- Jenkins is often chosen when there is a need for
  - Extreme customization
  - Supporting many SCM systems
  - Supporting complex legacy workflows
  - Infrastructure is managed and controlled
- GitHub / GitLab pipelines are often chosen when
  - Repositories are already hosted on the platform
  - Teams want fast setup
  - Minimal operational overhead is desired
  - Standard CI/CD patterns are sufficient

Which continuous integration (CI) tool do you regularly use in your company?




Image Credit: <https://blog.jetbrains.com/teamcity/2023/08/how-to-choose-cicd-tool/>

# Modern DevOps Ecosystems

- Jenkins is an automation engine that
  - Listens to SCM events
  - Executes pipelines
  - Orchestrates execution of builds, tests, scans, and deployments
  - Integrates with many external tools
    - *SCM: GitHub, GitLab, Bitbucket*
    - *Build tools: Maven, Gradle, npm, pip*
    - *Containers: Docker, Podman*
    - *Kubernetes: Jenkins Kubernetes plugin*
    - *Quality & Security: SonarQube, Snyk, Trivy*
    - *Secrets: Vault, cloud KMS systems*
- Jenkins is usually the best choice when
  - The pipeline needs to integrate across a variety of tool chains
  - Full control over pipeline workflows is required
  - Operational needs span hybrid or multi-cloud environments



# Jenkins Release Model

- Jenkins has two release lines
- Weekly releases
  - New features
  - Faster change
  - Higher risk
- LTS (Long-Term Support)
  - Stability-first
  - Security patches backported from the weekly releases
  - Recommended for production




Image Credit: <https://www.testingdocs.com/category/jenkins/>

# Jenkins Weekly Releases

- Represent the latest state of Jenkins development
- Typically used by
  - Plugin developers
  - Jenkins core contributors
  - Early adopters testing new features
  - Teams validating future LTS upgrades
- Not meant for
  - Production controllers
  - Regulated environments



# Jenkins LTS

- Recommended for production environments
  - Plugins are tested primarily against LTS
  - Predictable upgrade cycles
    - *Allows for enterprise planning for moving to newer versions*
  - Security patches are backported from the weekly build
    - *Ensures that security issues are addressed in LTS*
- Design assumes long running controller
  - Optimized for production environments



# Long Running Controller

- The controller
  - Maintains persistent state on disk
  - Accumulates configuration over time
- Retains
  - Job definitions
  - Build history
  - Security credentials
  - Plugin data
  - User configuration
- Different design architecture than
  - Stateless web services
  - Ephemeral CI runners
  - Serverless architectures




Image Credit: <https://naiwaen.debuggingsoft.com/2022/06/jenkins-controller-and-agents-architecture/>

# Long Running Controller

- Jenkins controllers
  - Retain long build histories
  - Preserve logs for auditing
  - Track trends over time
  - Examples
    - *Test flakiness trends*
    - *Build duration changes*
    - *Deployment frequency*
- This historical continuity only works if
  - The controller persists
  - The file system remains stable



# Jenkins Architecture

- Jenkins controller
  - Central brain of Jenkins
  - Responsibilities
    - *Web UI*
    - *REST API*
    - *Job and pipeline definitions*
    - *Build scheduling and queue management*
    - *Plugin management*
    - *Credential storage*
    - *User authentication and authorization*
    - *Should not run heavy builds*
  - Referred to as the master node in older versions of Jenkins




Image Credit: <https://naiwaen.debuggingsoft.com/2022/06/jenkins-controller-and-agents-architecture/>

# Jenkins Architecture

- Jenkins controller should NOT
  - Run heavy builds
  - Compile large projects
  - Execute tests
- Best practice is the builds are never executed on the controller node
  - Builds are done by agents
  - Jenkins controller is idle most of the time
  - But running builds might disrupt the scheduling processes that are the core of the controller functionality




Image Credit: <https://naiwaen.debuggingsoft.com/2022/06/jenkins-controller-and-agents-architecture/>

# Internal Controller Components

- Web application
  - Runs as a Java web app
  - Embedded Jetty servlet container
  - UI + API share the same backend
- Job configuration store
  - Jobs stored as XML on disk
  - Pipelines stored either:
    - *Inline (legacy)*
    - *From SCM (preferred)*
- Build queue
  - Pending builds waiting for executors
  - Queue logic decides
    - *When a job runs*
    - *On which agent*
    - *With which constraints*
- Executor management
  - Executors represent the capacity of an agent to run concurrent builds




Image Credit: <https://naiwaen.debuggingsoft.com/2022/06/jenkins-controller-and-agents-architecture/>

# Plugin Ecosystem

- Jenkins controller core is intentionally small
  - Most functionality comes from plugins
  - Thousands of plugins are available
  - Examples
    - *Git / GitHub / GitLab integration*
    - *Pipeline features*
    - *Kubernetes agents*
    - *Credentials and RBAC*
  - **Plugin Risks**
    - *Plugins run with high privileges*
    - *Outdated plugins = security risk*
    - *Version compatibility matters*




Image Credit: <https://thinkwithwp.com/blogs/devops/setting-up-a-ci-cd-pipeline-by-integrating-jenkins-with-aws-codebuild-and-aws-codedeploy/>

# Security Model

- Authentication: local users, LDAP, SSO
- Authorization: role-based access control (RBAC)
  - Credentials stored centrally and injected securely
  - Fine-grained permissions per job, folder, or pipeline
- Governance considerations
  - Who can install plugins?
  - Who can modify pipelines?
  - How are credentials managed?
  - How are updates tested and rolled out?




Image Credit: <https://thinkwithwp.com/blogs/devops/setting-up-a-ci-cd-pipeline-by-integrating-jenkins-with-aws-codebuild-and-aws-codedeploy/>

# Controller/Agent Architecture

- Using agents provides
  - Separation of concerns
    - *For example, doing builds in different OS environments*
  - Scalability – can add more agents to manage workloads
  - Security isolation – each agent can have its own security profile
- Static agents
  - Always running
  - Simple but resource-heavy
  - Common in older Jenkins setups
  - Usually a physical host or a virtual machine




Image Credit: <https://naiwaen.debuggingsoft.com/2022/06/jenkins-controller-and-agents-architecture/>

# Controller/Agent Architecture

- Ephemeral agents
  - Created on demand and destroyed after job completion
  - Often container-based or Kubernetes-based
  - Better isolation and scalability than static agents
  - Considered to be a modern best practice
- Execution flow of ephemeral agents
  - SCM event triggers pipeline
  - Controller schedules work to be done
  - Agent spins up
  - Pipeline stages execute the work
  - Agent is destroyed




Image Credit: <https://naiwaen.debuggingsoft.com/2022/06/jenkins-controller-and-agents-architecture/>

# SCM-Driven Automation

- Source control as the single source of truth
  - For both the application code and pipeline code (`Jenkinsfile`)
  - Pipelines ingest application code
  - `Jenkinsfile` is an IaC specification that defines the CI/CD steps and flow of execution
  - `Jenkinsfile` is versioned with the application
    - *Similar to how unit and integration tests are versioned*
  - Changes to `Jenkinsfile` code
    - *Follow code review processes*
    - *Subject to the same QA testing as application code*




Image Credit: <https://www.coachdevops.com/2020/06/how-to-configure-webhooks-in-bitbucket.html>

# SCM-Driven Automation

- Webhooks
  - GitHub/GitLab notify Jenkins on:
    - *Push*
    - *Pull/Merge Request*
    - *Tag creation*
  - Enables event-driven automation






Image Credit: <https://www.coachdevops.com/2020/06/how-to-configure-webhooks-in-bitbucket.html>

# SCM-Driven Automation

- Multibranch Pipelines
  - Jenkins automatically discovers branches
  - Each branch can have its own Jenkinsfile
- Enables
  - Pull request validation
  - Feature branch testing
  - Main branch deployments



# Installing Jenkins LTS

- Installing Jenkins means:
  - Installing the Jenkins controller
  - Running it as a long-lived service
  - Preparing it to manage jobs, pipelines, and agents
- Jenkins is a Java application distributed as
  - Native packages (RPM, DEB) or Windows MSI
  - WAR file
  - Docker image
- Always start with LTS unless there is a reason for using the weekly build



# Installing Jenkins LTS

- Native Package Install
  - Installed as a system service
  - Runs on a VM or server
  - Direct access to filesystem
  - Traditional enterprise model
- WAR File
  - Portable
  - Manual startup
  - Often used for testing or learning
- Docker
  - Containerized Jenkins controller



# Installing Jenkins LTS

- During the first startup
  - Jenkins creates \$JENKINS\_HOME
  - Default configuration files are generated
  - Initial admin password is created
  - Setup wizard is triggered

Getting Started

## Unlock Jenkins

To ensure Jenkins is securely set up by the administrator, a password has been written to the log (not sure where to find it?) and this file on the server:

`/var/jenkins_home/secrets/initialAdminPassword`

Please copy the password from either location and paste it below.

Administrator password

Continue



# Initial Setup Wizard

- On first startup, Jenkins
  - Locks itself
  - Requests an admin password
  - Prompts for plugin installation
  - Creates the first admin user
  - Sets basic configuration

Getting Started

## Unlock Jenkins

To ensure Jenkins is securely set up by the administrator, a password has been written to the log (not sure where to find it?) and this file on the server:

`/var/jenkins_home/secrets/initialAdminPassword`

Please copy the password from either location and paste it below.

Administrator password

Continue

# Plugin Selection

- Two common choices:
  - Suggested plugins
  - Custom plugin selection
- For most installations
  - Suggested plugins are sufficient
  - Plugins can always be added later
  - However, removing them is harder

The screenshot shows the 'Available plugins' section of the Blue Ocean plugin management interface. The search bar at the top contains the text 'blue ocean'. Below the search bar, there are four tabs: 'Install', 'Name ↓', 'Released', and 'Last updated ↑'. The 'Install' tab is selected. The results list includes the following items:

Plugin Name	Description	Last Updated
Blue Ocean 1.27.3	External Site/Tool Integrations   User Interface BlueOcean Aggregator	1 day 2 hr ago
Display URL for Blue Ocean 2.4.1	This plugin generates BlueOcean specific URLs for the Display URL plugin.	2 yr 1 mo ago
Personalization for Blue Ocean 1.27.3	External Site/Tool Integrations   User Interface Blue Ocean Personalization	1 day 2 hr ago
Bitbucket Pipeline for Blue Ocean 1.27.3	BlueOcean Bitbucket pipeline creator	1 day 2 hr ago

At the bottom of the interface, there are two buttons: 'Install without restart' and 'Download now and install after restart'. To the right of these buttons, the text 'Update information obtained: 1 hr 18 min ago' and 'Check now' are displayed.

# Global Tools Configuration

- Global tools:
  - Are *not* binaries
  - Are *not* installed on the controller
  - Are *not* automatically present on agents
  - Are named tool definitions
  - Are stored in Jenkins configuration
- Used by pipelines as references
  - Global tools are labels, not software

The screenshot shows the Jenkins 'Tools' configuration page. At the top, there is a navigation bar with the Jenkins logo, the text 'Jenkins / Manage Jenkins / Tools', and three icons: a magnifying glass (search), a gear (settings), and a person (user). Below the navigation, the title 'Tools' is displayed, followed by the subtitle 'Configure tools, their locations and automatic installers.' Under the 'Maven Configuration' section, there is a 'Default settings provider' dropdown set to 'Use default maven settings'. Below it, a 'Default global settings provider' dropdown is also set to 'Use default maven global settings'. In the 'JDK installations' section, there is a single entry labeled 'Use default jdk installation'. The page has a light blue header and a white main content area.

# Reference Tools by Name

- In a Jenkinsfile:

- “Use Maven version maven-3.9”
- The name “maven-3.9” is defined globally in Jenkins
- Abstracts away
  - *Version numbers*
  - *Installation paths*
  - *Agent OS differences*

- Benefits

- Pipelines become portable
- Jenkins controls versions centrally
- Developers don’t hardcode paths to build tools in pipeline code

The screenshot shows the Jenkins interface for managing tools. At the top, there's a navigation bar with icons for search, settings, and user profile. The main title is "Jenkins / Manage Jenkins / Tools". Below the title, the section "Tools" is displayed with the sub-instruction "Configure tools, their locations and automatic installers." Under "Maven Configuration", there are two dropdown menus: "Default settings provider" set to "Use default maven settings" and "Default global settings provider" also set to "Use default maven global settings". At the bottom, there's a section titled "JDK installations" which is currently empty.

# What Happens at Runtime

- Step 1: Pipeline starts
  - Controller reads the Jenkinsfile
  - Sees: "This pipeline needs Maven maven-3.9"
- Step 2: Agent is selected
  - Controller assigns an agent
  - Agent may be:
    - A VM
    - A Docker container
    - A Kubernetes pod
- Step 3: Tool resolution
  - Controller sends the tool definition to the agent
  - Definition includes:
    - *Tool type (Maven, JDK, Node)*
    - *Version*
    - *Installation strategy*

## Maven installations

Maven installations ^ Edited

+ Add Maven


**Maven**

Name: Maven-3.9.11

MAVEN\_HOME: /usr/share/maven

Install automatically ?

+ Add Maven



# What Happens at Runtime

- Step 4: Tool installation (if needed)
  - On the agent, if tool is already installed then it is used
  - If not
    - Jenkins downloads it
    - Installs it locally
    - Caches it for reuse (if agent persists)
- Step 5: Tool execution
  - Agent updates PATH and environment variables
  - Pipeline steps run
  - Controller receives logs and results

## Maven installations

The screenshot shows the 'Maven installations' configuration screen in Jenkins. At the top, there's a header with 'Maven installations' and an 'Edited' status indicator. Below the header is a button labeled '+ Add Maven'. The main area contains a single Maven entry with the following fields:

- Name:** Maven-3.9.11
- MAVEN\_HOME:** /usr/share/maven
- Install automatically:** An unchecked checkbox with a question mark icon.

At the bottom of the list is another '+ Add Maven' button.



# What Auto-Install Means

- Jenkins can download tools automatically
- Controlled by:
  - Tool definition settings
  - Plugin support
  - Agent permissions
- Benefits
  - No manual agent setup
  - New agents work immediately
  - Ephemeral agents become practical
- Pitfalls
  - Builds depend on external downloads
  - Tool versions can change unexpectedly
  - Network failures can break builds

Maven installations

Maven installations ^ Edited

+ Add Maven


Maven

Name: Maven-3.9.11

MAVEN\_HOME: /usr/share/maven

Install automatically ?

+ Add Maven



# Ensuring Consistent Builds

- Without global tools
  - Each agent might have different versions
  - Pipelines hardcode paths
  - “Works on my agent” problems appear
- With global tools
  - Jenkins enforces version consistency
  - Pipelines are environment-agnostic
  - Upgrades are centralized and controlled
- Critical when
  - Multiple teams share Jenkins
  - Agents are ephemeral
  - Jenkins uses containers or Kubernetes

```
pipeline {  
    agent any  
  
    tools {  
        // These refer to Global Tool Configuration names  
        jdk 'JDK17'  
        maven 'Maven3'  
    }  
  
    stages {  
        stage('Checkout') {  
            steps {  
                git 'https://github.com/example/my-java-app.git'  
            }  
        }  
  
        stage('Build') {  
            steps {  
                sh 'mvn clean package'  
            }  
        }  
    }  
}
```



# Controller vs Agent Tooling


- Controller responsibilities
  - Stores tool definitions
  - Defines names and versions
  - Decides what tools are needed
  - Never runs builds
- Agent responsibilities
  - Installs tools
  - Executes commands
  - Runs builds and tests
  - Cleans up afterward
- The controller is the planner
- The agent is the worker

```
pipeline {  
    agent any  
  
    tools {  
        // These refer to Global Tool Configuration names  
        jdk 'JDK17'  
        maven 'Maven3'  
    }  
  
    stages {  
        stage('Checkout') {  
            steps {  
                git 'https://github.com/example/my-java-app.git'  
            }  
        }  
  
        stage('Build') {  
            steps {  
                sh 'mvn clean package'  
            }  
        }  
    }  
}
```




# Running Jenkins in Docker

- Docker helps solve common Jenkins problems
  - Consistent environments
  - Easier upgrades
  - Simplified rollback
  - Faster setup for labs and training
- Jenkins in Docker is NOT Jenkins as a stateless service
  - The controller still needs persistent storage




# Running Jenkins in Docker

- Persistent volumes
  - Jenkins state must survive container restarts
  - `$JENKINS_HOME` is mounted to a volume
  - Losing this volume = losing Jenkins
- Container Lifecycle
  - Container can be destroyed and recreated
  - Jenkins data remains intact
  - Encourages clean operational practices




# Running Jenkins in Docker

- Running Jenkins in Docker means
  - Jenkins runs inside a container
  - Jenkins data lives in a mounted volume
  - The container can be replaced without losing Jenkins state
  - Provides control, consistency, and safety without changing Jenkins' architecture




# Docker: Fast and Predictable Setup

- Without Docker
  - Install Java
  - Configure OS packages
  - Manage service startup
  - Troubleshoot environment differences
- With Docker
  - Pull the Jenkins LTS image
  - Start a container
  - Jenkins is immediately usable
  - Everyone gets the same Jenkins
  - Avoids setup chaos




# Docker: Environment Consistency

- Docker ensures
  - Same Jenkins version
  - Same Java version
  - Same default configuration
  - Same filesystem layout
- Eliminates
  - “Works on my machine”
  - OS-specific surprises
  - Hidden dependencies




# Docker: Modern Jenkins Practices

- Jenkins has evolved to support
  - Configuration-as-Code (JCasC)
  - Ephemeral agents
  - Cloud deployments
  - Kubernetes deployments




# Docker: Reduced Operational Risk

- Avoids common Jenkins failures
  - Broken Java upgrades
  - OS package conflicts
  - Accidental system changes
  - Hard-to-reproduce environments
- Important because
  - Jenkins controllers are long-running
  - Downtime is disruptive
  - Recovery needs to be predictable



# When to Avoid Docker

- Very high-scale enterprise setups
- Strict compliance environments
- Existing Jenkins infrastructure already standardized



# Configuration-as-Code (JCasC)

- Configuration-as-Code (JCasC) is where
  - Jenkins configurations are defined in YAML files
  - Loads the configuration automatically at startup
  - Treats Jenkins configuration like application code
- Declarative
  - Describes how Jenkins should be configured
  - Replaces configuring Jenkins through the UI

```
jenkins:  
    systemMessage: "Welcome to the Jenkins Training Instance"  
  
    securityRealm:  
        local:  
            allowsSignup: false  
        users:  
            - id: admin  
                password: admin123  
  
    authorizationStrategy:  
        loggedInUsersCanDoAnything:  
            allowAnonymousRead: false  
  
tool:  
    git:  
        installations:  
            - name: Default Git  
                home: /usr/bin/git
```

# Configuration-as-Code (JCasC)

- JCasC, can define
  - Security settings
  - Authentication and authorization
  - Users and roles
  - Credentials
  - Global tools (JDK, Maven, Node, etc.)
  - Plugin configuration
  - System settings
- But not
  - Individual pipeline logic
  - Build steps inside Jenkinsfiles
  - Pipelines are code, JCasC is configuration.

```
jenkins:  
    systemMessage: "Welcome to the Jenkins Training Instance"  
  
    securityRealm:  
        local:  
            allowsSignup: false  
            users:  
                - id: admin  
                  password: admin123  
  
    authorizationStrategy:  
        loggedInUsersCanDoAnything:  
            allowAnonymousRead: false  
  
tool:  
    git:  
        installations:  
            - name: Default Git  
              home: /usr/bin/git
```



# Configuration-as-Code (JCasC)

- JCasC YAML files are stored on disk
- Often mounted into Jenkins via:
  - Docker volumes
  - ConfigMaps (Kubernetes)
- JCasC is applied at
  - Jenkins startup
  - Configuration reload (optional)
- What Jenkins does
  - Jenkins starts
  - JCasC plugin reads YAML
  - Jenkins configures itself accordingly
  - UI reflects the defined state

```
jenkins:  
    systemMessage: "Welcome to the Jenkins Training Instance"  
  
    securityRealm:  
        local:  
            allowsSignup: false  
        users:  
            - id: admin  
                password: admin123  
  
    authorizationStrategy:  
        loggedInUsersCanDoAnything:  
            allowAnonymousRead: false  
  
tool:  
    git:  
        installations:  
            - name: Default Git  
                home: /usr/bin/git
```



# Benefits

- Makes Jenkins reproducible
- Supports automation and version control
- Enables safer upgrades
- Helps avoid snowflake Jenkins which
  - Were configured manually over time
  - Have unique, undocumented settings
  - Cannot be easily recreated
  - No one fully understands anymore
  - Breaks when you try to upgrade or migrate
- Called a snowflake because it's unique and melts under pressure

```
jenkins:  
    systemMessage: "Welcome to the Jenkins Training Instance"  
  
    securityRealm:  
        local:  
            allowsSignup: false  
        users:  
            - id: admin  
                password: admin123  
  
    authorizationStrategy:  
        loggedInUsersCanDoAnything:  
            allowAnonymousRead: false  
  
tool:  
    git:  
        installations:  
            - name: Default Git  
                home: /usr/bin/git
```



# Snowflake Instances

- Common causes

- Clicking through the UI to configure Jenkins
- Installing plugins “just to try something”
- Making emergency fixes directly on the server
- Manually editing config files
- Installing tools directly on agents
- No version control for configuration

- Over time

- The original admins leave
- Configuration knowledge is lost

```
jenkins:  
    systemMessage: "Welcome to the Jenkins Training Instance"  
  
    securityRealm:  
        local:  
            allowsSignup: false  
        users:  
            - id: admin  
                password: admin123  
  
    authorizationStrategy:  
        loggedInUsersCanDoAnything:  
            allowAnonymousRead: false  
  
tool:  
    git:  
        installations:  
            - name: Default Git  
                home: /usr/bin/git
```

# Questions

